

Hadoop Map Reduce and YARN

Lecture 4

Tahmina Sultana Priya
Lecturer, Dept.of CSE
Daffodil International University

Agenda for today's Session

edure

1. What is Hadoop MapReduce?
2. MapReduce In Nutshell
3. Advantages of MapReduce
4. Hadoop MapReduce Approach with an Example
5. Hadoop MapReduce/YARN Components
6. YARN With MapReduce
7. Yarn Application Workflow
8. MapReduce Program with Hands On



2 main Hadoop Components

Storage

Processing



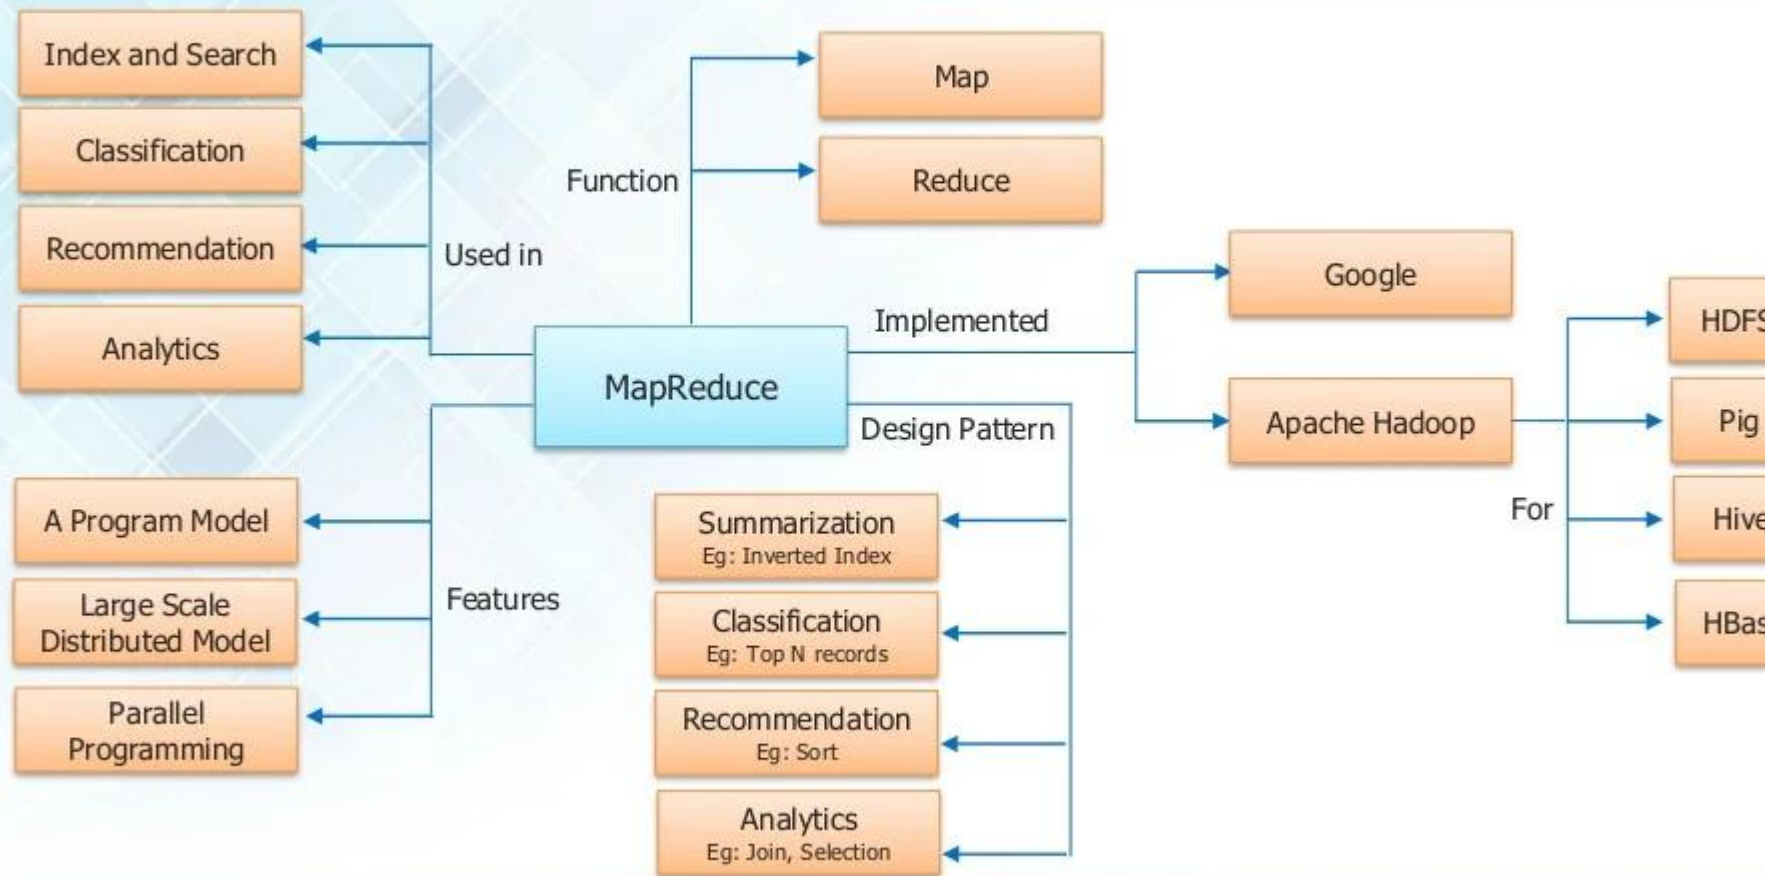
MapReduce: Data Processing Using Programming

edure



- *Hadoop MapReduce is the processing component of Apache Hadoop*
- *It processes data parallelly in distributed environment*



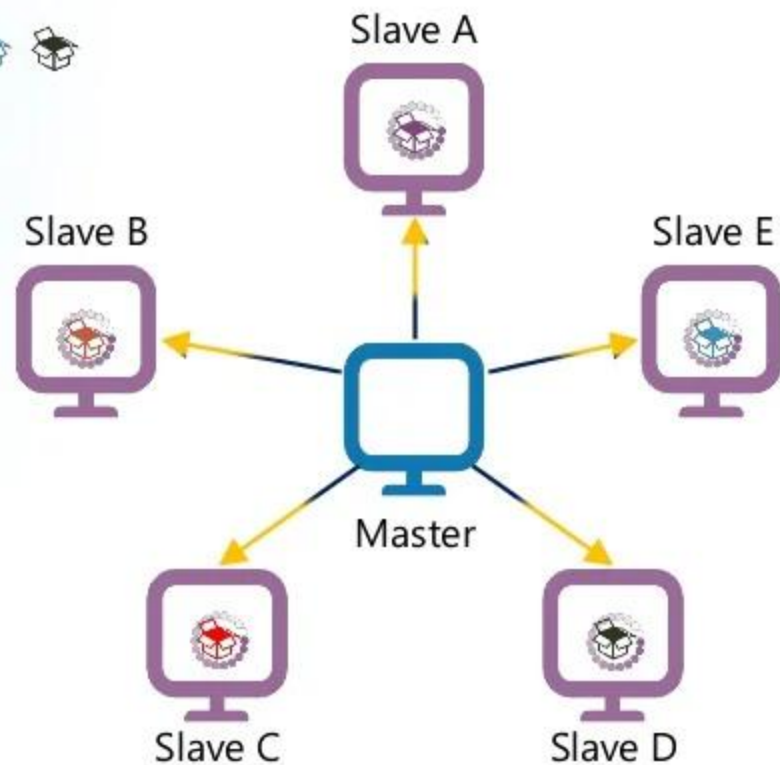


2 Biggest Advantages of MapReduce

Advantage 1: Parallel Processing

Data → 

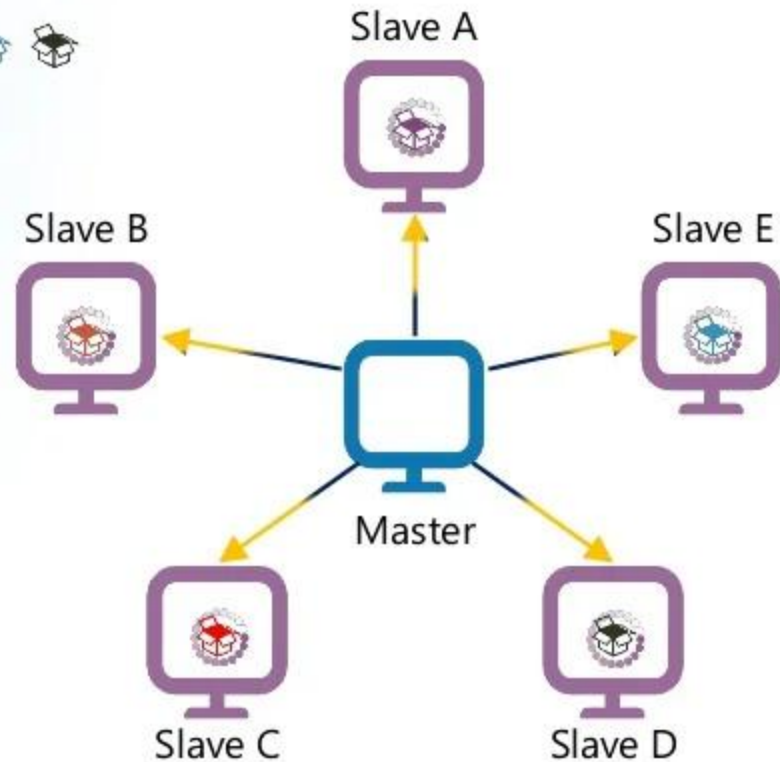
- Data is processed in parallel
- Processing becomes fast



Advantage 2: Data Locality - Processing to Storage edure

Data → 

- Moving Data to processing is very costly
- In MapReduce, we move processing to Data

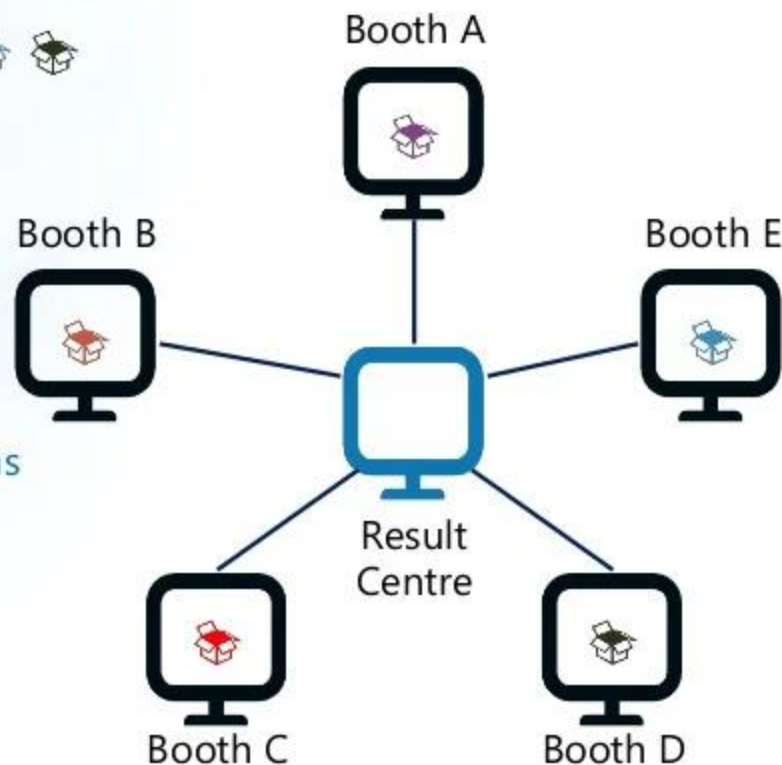


Traditional vs MapReduce Way

Data →     

Election Votes Casting


- Votes is stored at different Booths
- Result Centre has the details of all the Booths



Data → 

Counting – Traditional Approach

- Votes are moved to Result Centre for counting
- Moving all the votes to Centre is costly
- Result Centre is over-burdened
- Counting takes time




Hadoop MapReduce To the Rescue!

edure

Data → 

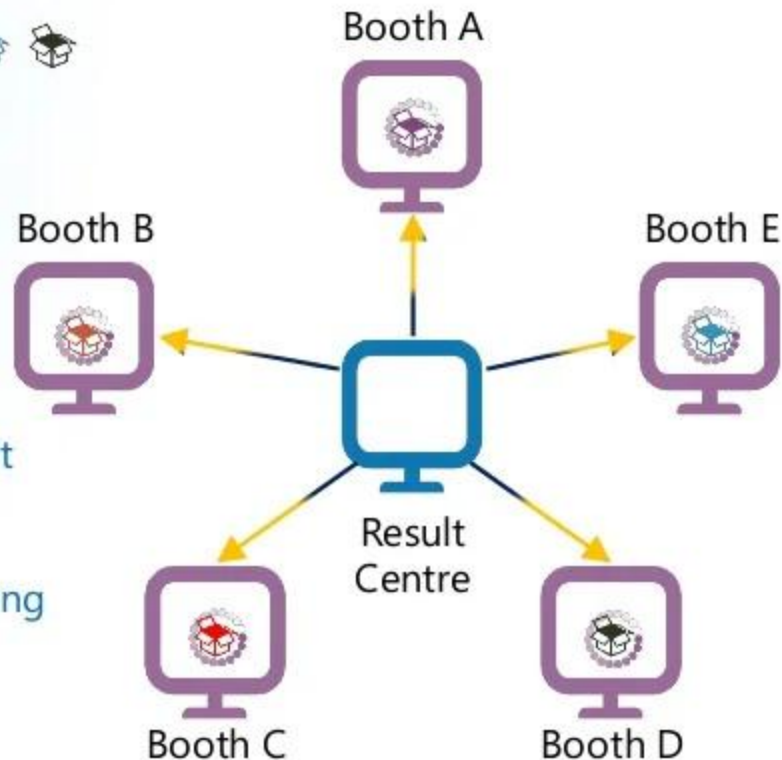
✗ Hadoop MapReduce Doesn't Follow This Approach



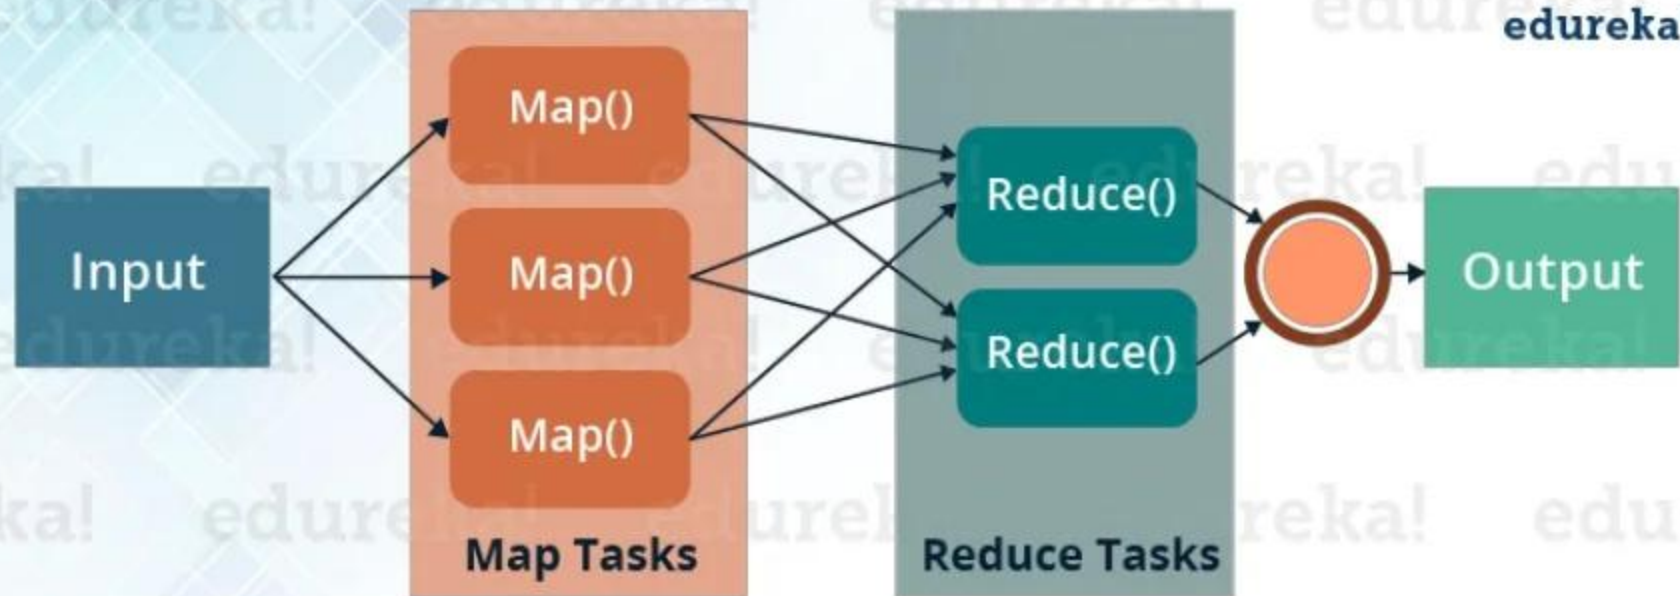
Votes → 

Counting – MapReduce Approach

- Votes are counted at individual booths
- Booth-wise results are sent back to the result centre
- Final Result is declared easily and quickly using this way




MapReduce In Detail



Anatomy of a MapReduce Program

edure



Map:

(K1, V1)

List (K2, V2)

Reduce:

(K2, list (V2))

List (K3, V3)

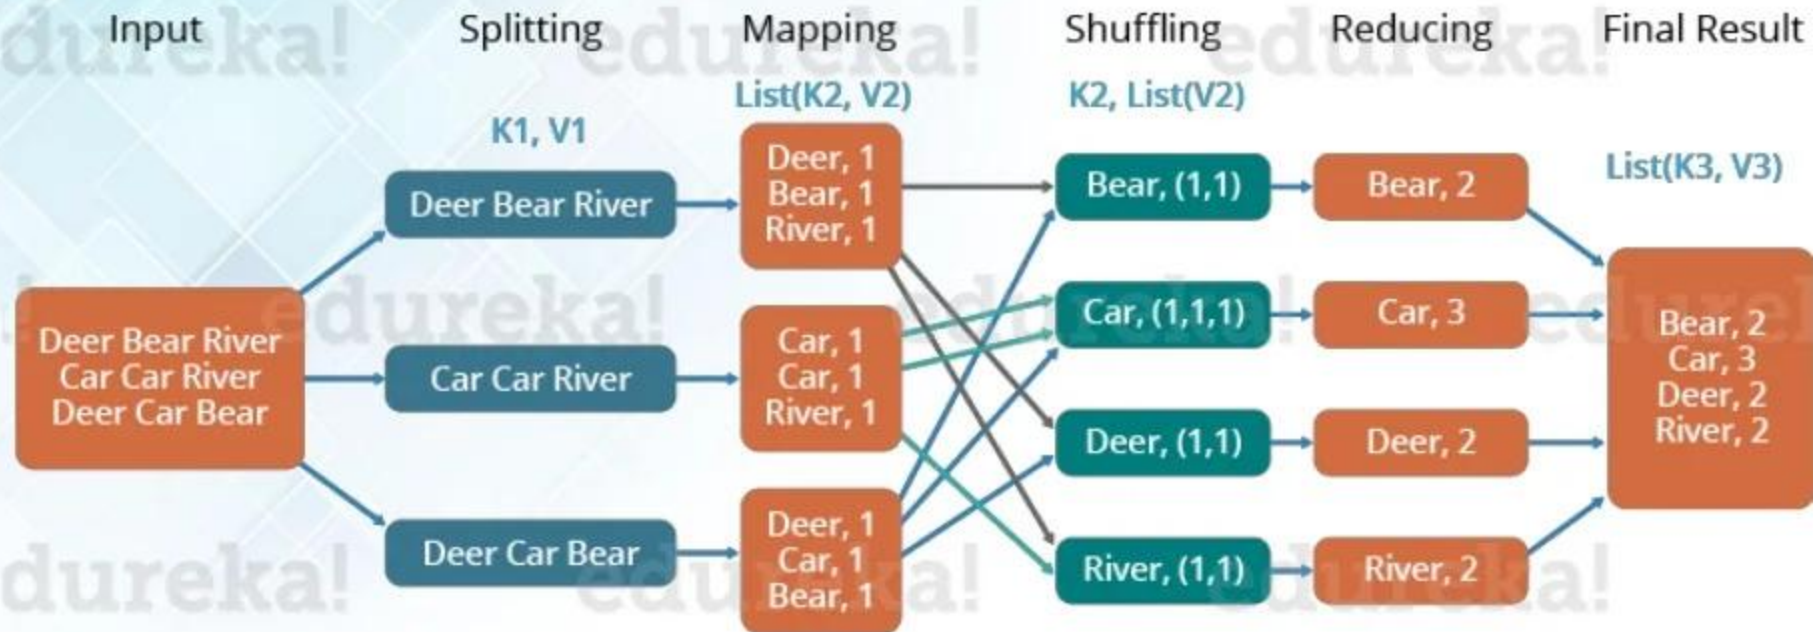
Let us take an example to understand
MapReduce Way

MapReduce Way – Word Count Process

edureka!

The Overall MapReduce Word Count Process

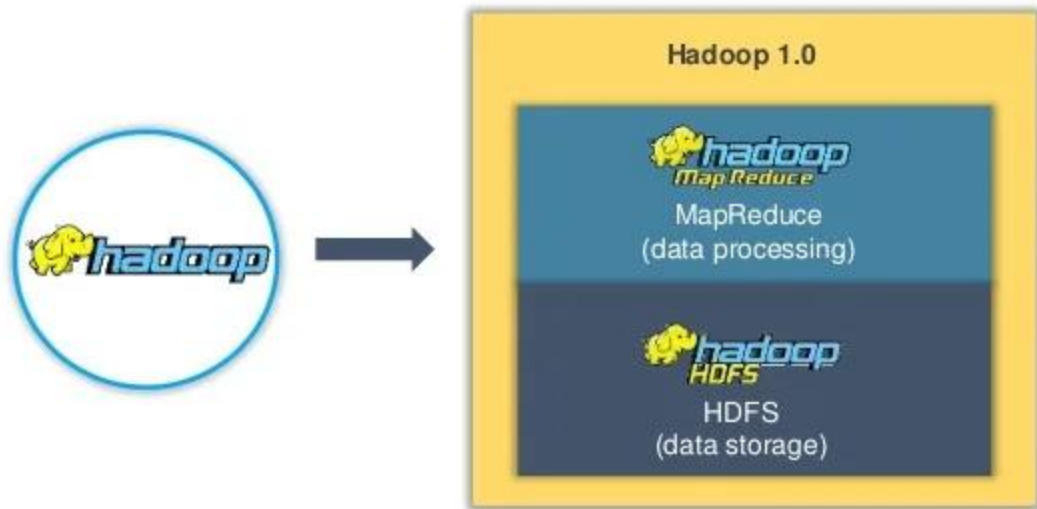
edureka!



Hadoop 1.0 (MR 1)




Hadoop 1.0 (MR 1)



In Hadoop 1.0, MapReduce performed both data processing and resource management



Data processing



Resource management

Hadoop 1.0 (MR 1)



MapReduce consisted of
Job Tracker and Task Tracker



Job
Tracker

Allocated resources, performed
scheduling and monitored jobs

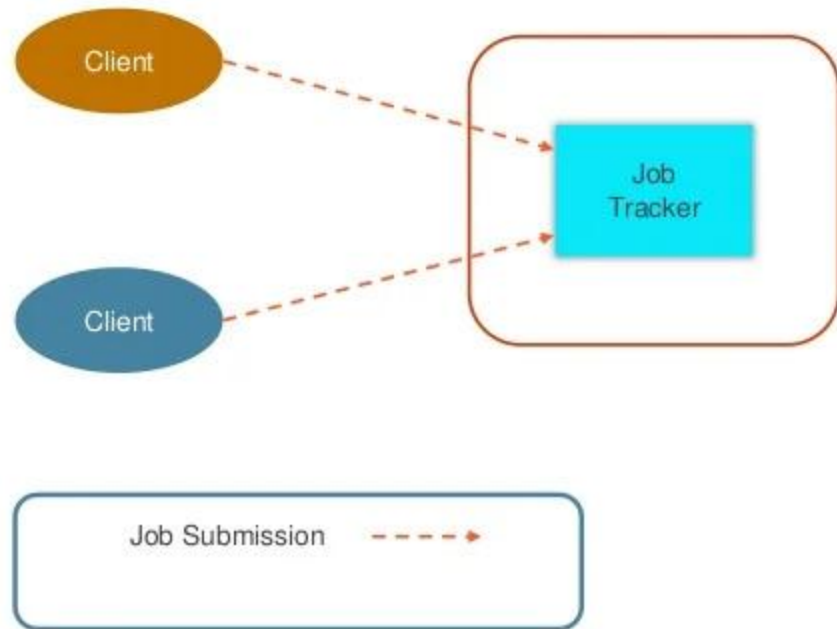
Assigned map and reduce tasks to jobs
running on Task Trackers

Task
Tracker

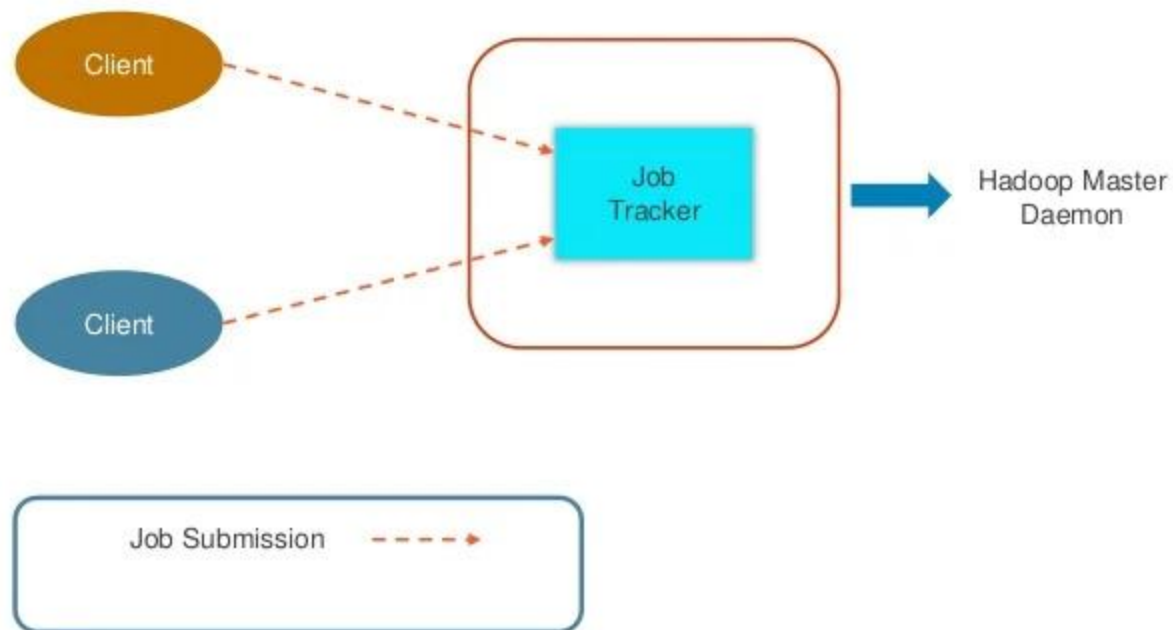
Task Trackers processed the jobs

Task Trackers reported their progress
to the Job Tracker





Hadoop 1.0 (MR 1)



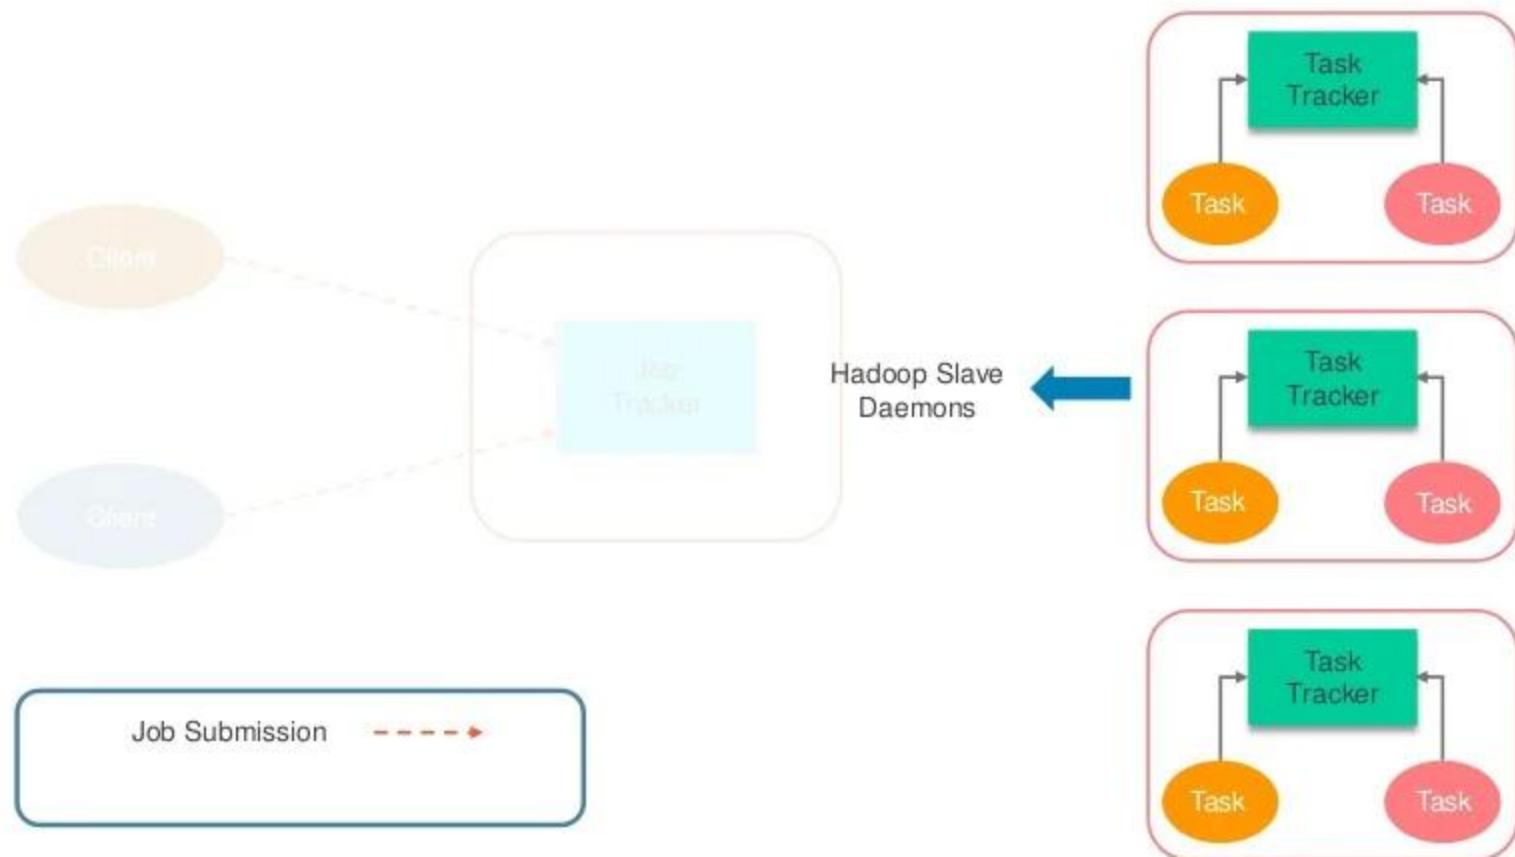
Hadoop 1.0 (MR 1)



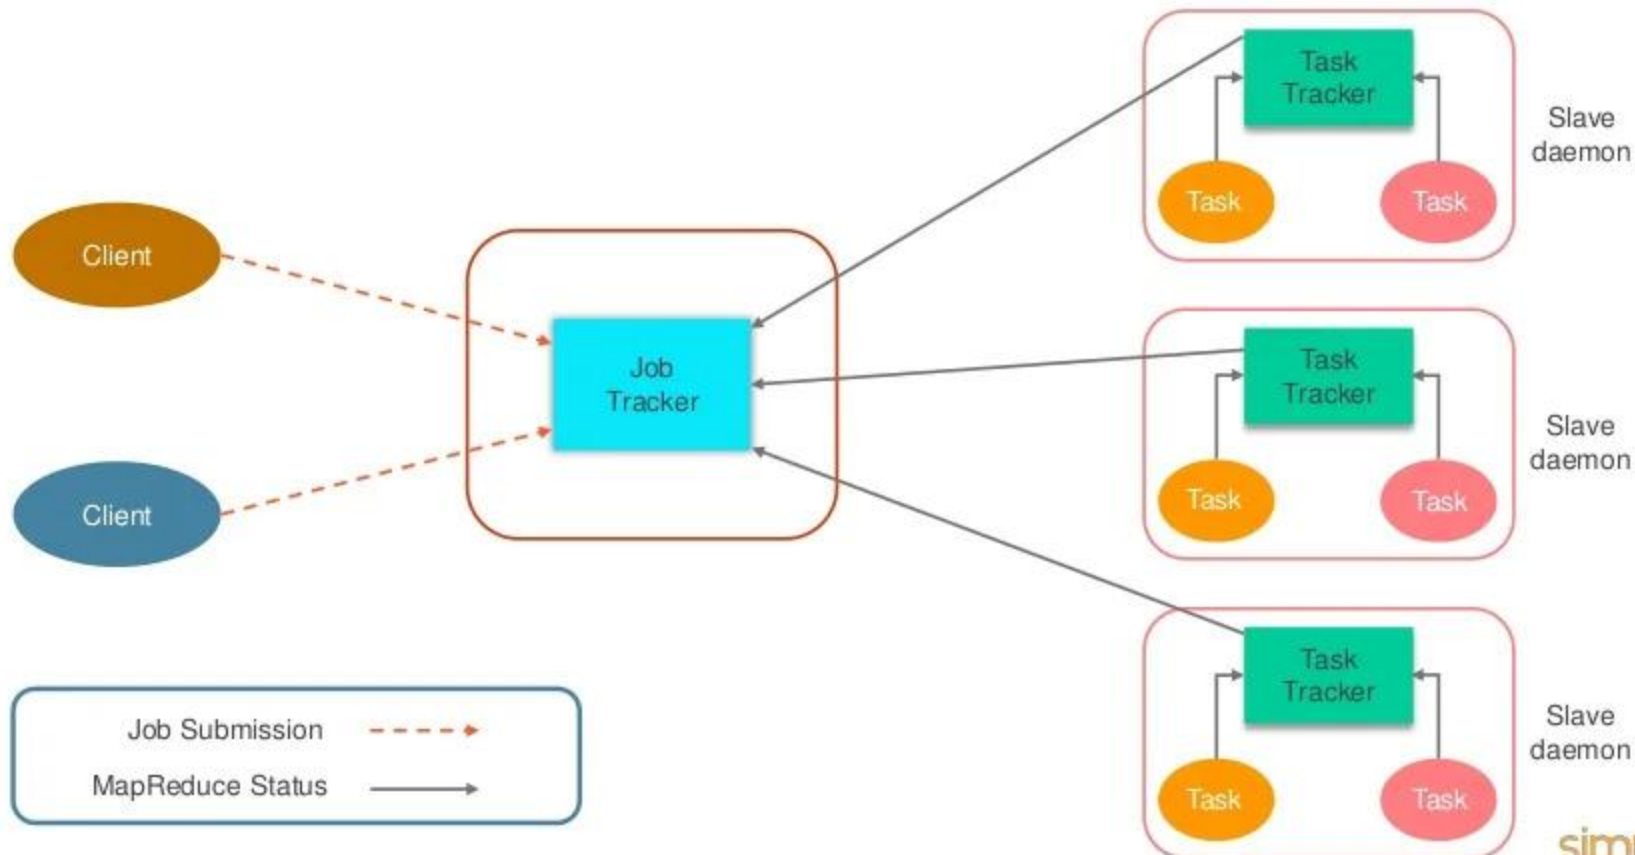
Hadoop 1.0 (MR 1)



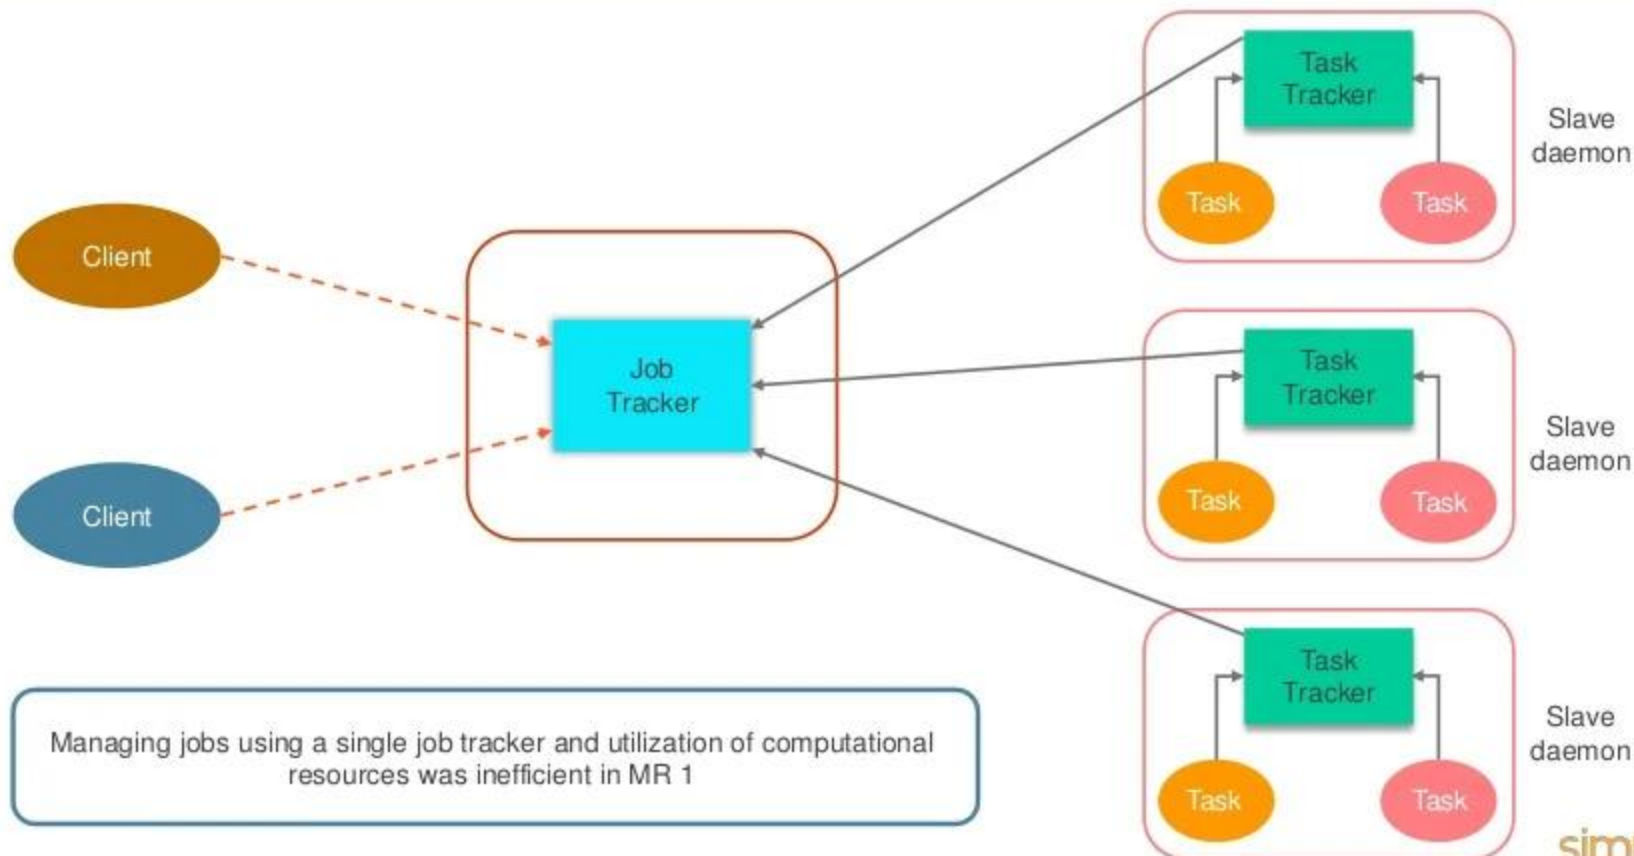
Hadoop 1.0 (MR 1)



Hadoop 1.0 (MR 1)



Hadoop 1.0 (MR 1)



Limitations of Hadoop 1.0 (MR 1)

1

Scalability

Due to a [single JobTracker](#), scalability became a bottleneck.

Cannot have a cluster size of more than [4000 nodes](#) and cannot run more than [40000 concurrent tasks](#)



Limitations of Hadoop 1.0 (MR 1)

1

Scalability

Due to a single JobTracker, scalability became a bottleneck.

Maximum cluster size - 4000 nodes
Maximum concurrent tasks - 40000

2

Availability issue

JobTracker is **single point of failure**. Any failure kills all queued and running jobs. Jobs need to be resubmitted by users

Limitations of Hadoop 1.0 (MR 1)

3

Resource Utilization

Due to predefined number of **map**
and **reduce slots** for each
TaskTracker, **resource utilization**
issues occur



Limitations of Hadoop 1.0 (MR 1)

3

Resource Utilization

Due to pre-defined number of map and reduce slots for each TaskTracker, resource utilization issues occur

4

Limitations in running non-MapReduce applications

Problem in performing **real-time analysis** and running **Ad-hoc query** as MapReduce is batch driven

MapReduce Using Yarn


Need for YARN



Need for YARN


Before YARN

Hadoop 1.0

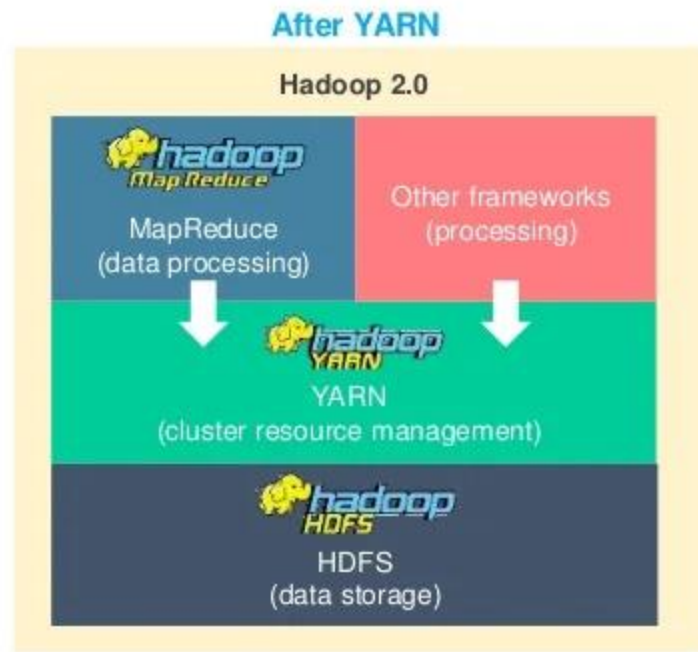


Designed to run MapReduce jobs only and had issues in scalability, resource utilization, etc.

Need for YARN



Designed to run MapReduce jobs only and had issues in scalability, resource utilization, etc.



YARN solved those issues and users could work on multiple processing models along with MapReduce

Solution - Hadoop 2.0 (YARN)



Scalability



Can have a cluster size of more than 10,000 nodes and can run more than 1,00,000 concurrent tasks

Solution - Hadoop 2.0 (YARN)



Scalability



Can have a cluster size of more than 10,000 nodes and can run more than 1,00,000 concurrent tasks

Compatibility



Applications developed for Hadoop 1 runs on YARN without any disruption or availability issues

Solution - Hadoop 2.0 (YARN)



Scalability




Can have a cluster size of more than 10,000 nodes and can run more than 1,00,000 concurrent tasks

Compatibility



Applications developed for Hadoop 1 runs on YARN without any disruption or availability issues

Resource utilization



Allows dynamic allocation of cluster resources to improve resource utilization

Solution - Hadoop 2.0 (YARN)



Scalability




Can have a cluster size of more than 10,000 nodes and can run more than 1,00,000 concurrent tasks

Compatibility



Applications developed for Hadoop 1 runs on YARN without any disruption or availability issues

Resource utilization



Allows dynamic allocation of cluster resources to improve resource utilization

Multitenancy



Can use open-source and propriety data access engines and perform real-time analysis and running ad-hoc query

YARN – Moving beyond MapReduce

edure



*Applications Run Natively **IN** Hadoop*

BATCH
(MapReduce)

INTERACTIVE
(Text)

ONLINE
(HBase)

STREAMING
(Storm, S4, ...)

GRAPH
(Giraph)

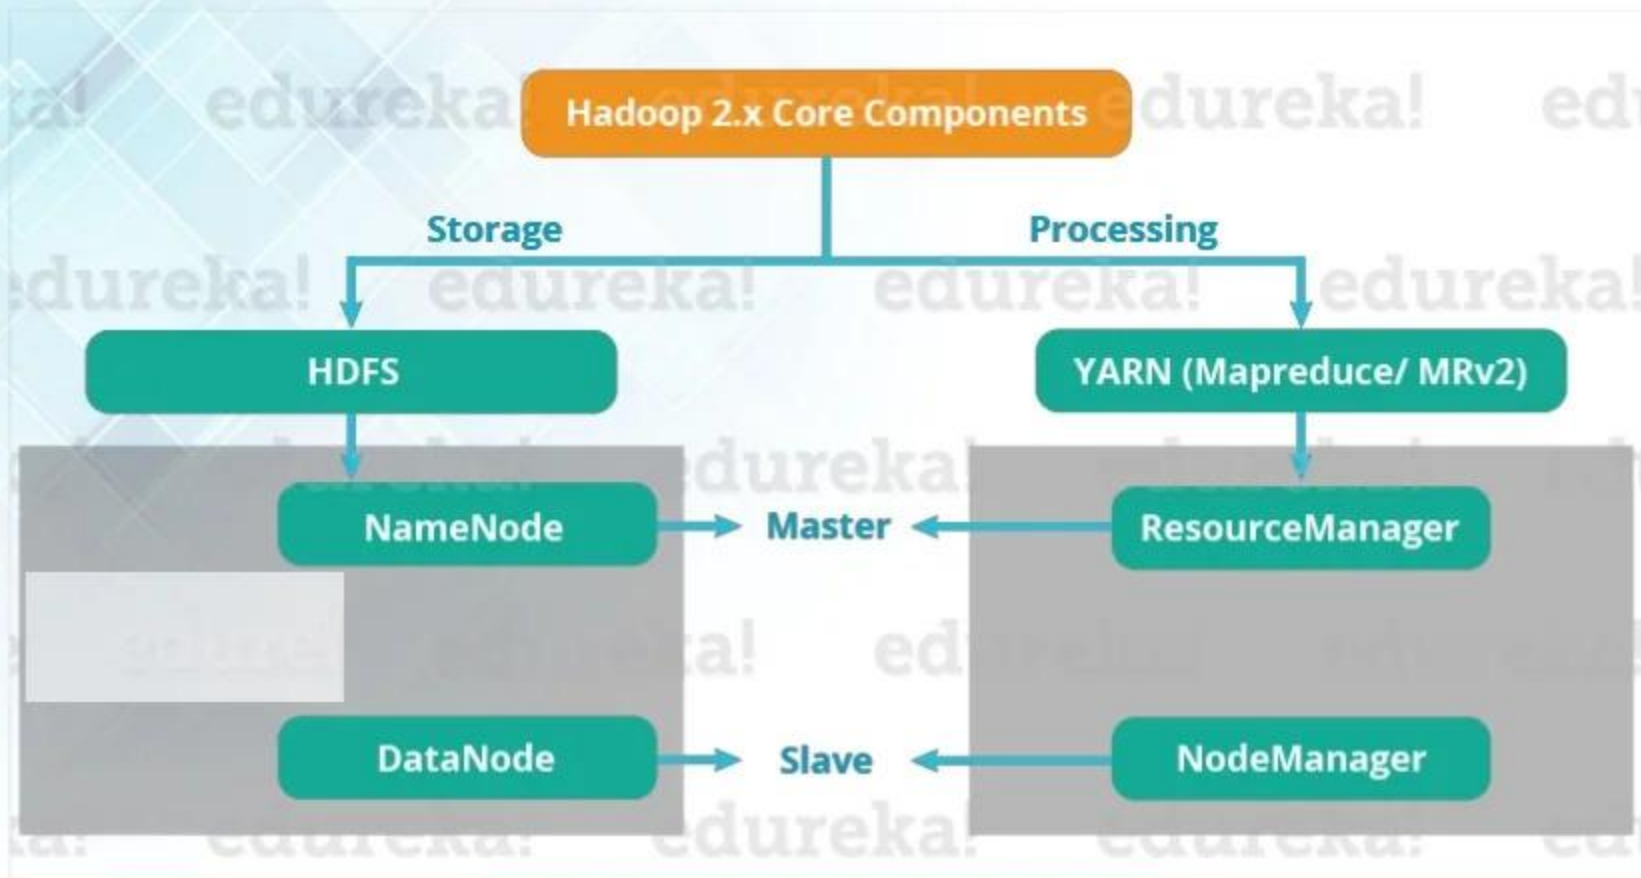
IN-MEMORY
(Spark)

HPC MPI
(OpenMPI)

OTHER
(Search)
(Weave..)

YARN (Cluster Resource Management)

HDFS2 (Redundant, Reliable Storage)



→ Client

- » Submits a MapReduce Job

→ Resource Manager

- » Cluster Level resource manager
- » Long Life, High Quality Hardware

→ Node Manager

- » One per Data Node
- » Monitors resources on Data Node

→ Job History Server

- » Maintains information about submitted MapReduce jobs after their ApplicationMaster terminates

→ ApplicationMaster

- » One per application
- » Short life
- » Coordinates and Manages MapReduce Jobs
- » Negotiates with Resource Manager to schedule tasks
- » The tasks are started by NodeManager(s)

→ Container

- » Created by NM when requested
- » Allocates certain amount of resources (memory, CPU etc.) on a slave node


YARN Architecture







YARN Architecture




YARN Architecture




YARN Architecture




YARN Architecture



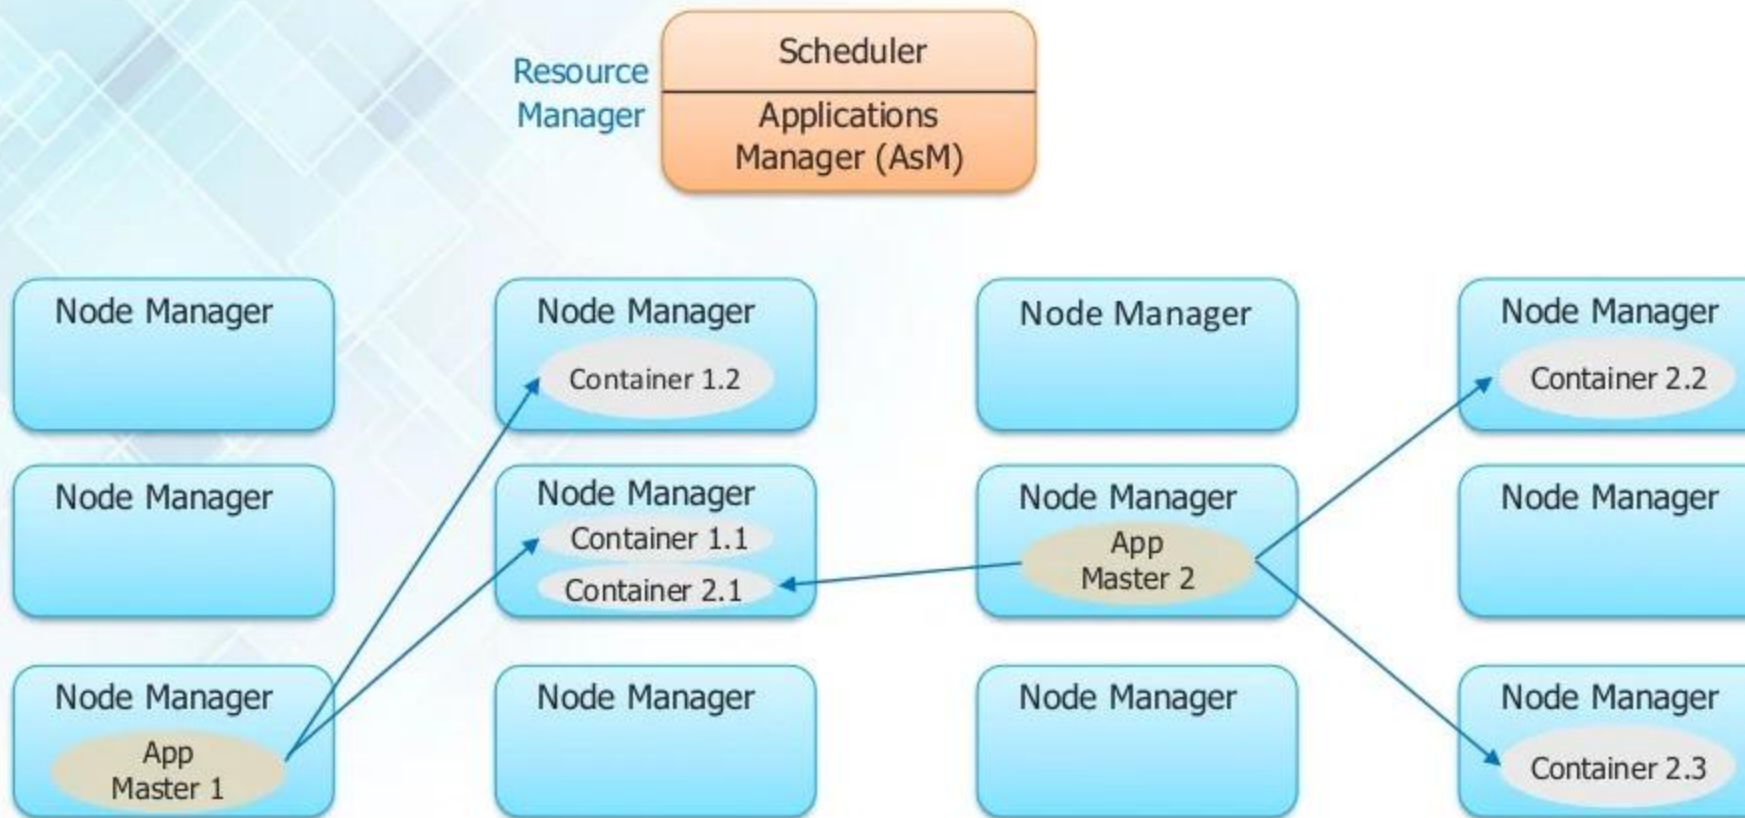
YARN Architecture



YARN Architecture




YARN Application Workflow in MapReduce




→ Execution Sequence :

1. Client submits an application




→ Execution Sequence :

1. Client submits an application
2. RM allocates a container to start AM




→ Execution Sequence :

1. Client submits an application
2. RM allocates a container to start AM
3. AM registers with RM




→ Execution Sequence :

1. Client submits an application
2. RM allocates a container to start AM
3. AM registers with RM
4. AM asks containers from RM



→ Execution Sequence :


1. Client submits an application
2. RM allocates a container to start AM
3. AM registers with RM
4. AM asks containers from RM
5. AM notifies NM to launch containers



Application Workflow

→ Execution Sequence :


1. Client submits an application
2. RM allocates a container to start AM
3. AM registers with RM
4. AM asks containers from RM
5. AM notifies NM to launch containers
6. Application code is executed in container



Application Workflow


→ Execution Sequence :

1. Client submits an application
2. RM allocates a container to start AM
3. AM registers with RM
4. AM asks containers from RM
5. AM notifies NM to launch containers
6. Application code is executed in container
7. Client contacts RM/AM to monitor application's status



→ Execution Sequence :

1. Client submits an application
2. RM allocates a container to start AM
3. AM registers with RM
4. AM asks containers from RM
5. AM notifies NM to launch containers
6. Application code is executed in container
7. Client contacts RM/AM to monitor application's status
8. AM unregisters with RM



Thank You

Slide content from
Edureka and Simplilearn